

IDPT-FP (IDPT, Full Protocol)

Jorge García Vidal, UPC, BSC

Draft version May 23th 2020

This is a proposal for a digital proximity tracing protocol that can operate both in centralized and distributed mode with full interoperability. It is based on the mechanism described for the IDPT protocol¹ for interoperability between ROBERT² and DP3T³ applications.

Assumptions

We assume that in the same geographical area (e.g. the Schengen area) we have a digital proximity tracking application with users who can decide whether the risk score is done on a central server (C-users), or is done on the user's phones (D-users).

It may be that in a country within this geographical area, national public health authorities choose to support only C-users or D-users. Another option is that they give freedom of choice to their citizens, who assess the trade-offs between privacy and security and the effectiveness of the risk score.

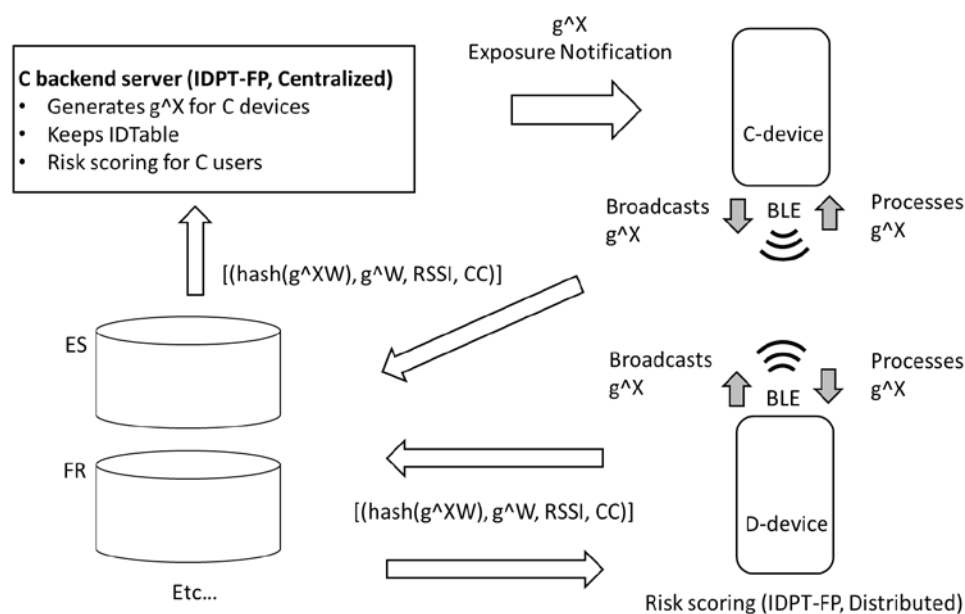


Figure: Functional elements

¹ <https://github.com/IDPTdocs/documents/blob/master/IDPT-v2.pdf>

² https://github.com/ROBERT-proximity-tracing/documents/blob/master/ROBERT-specification-EN-v1_0.pdf

³ <https://github.com/DP-3T/documents/blob/master/DP3T%20White%20Paper.pdf>

Beacon broadcast

- All users broadcast BLE ADV_IND packets (which we call "beacons") with a payload equal to g^X , as in a Diffie-Hellman exchange, where X is a secret number that is changed at each epoch (e.g. 15 minutes).
- We assume that g^X is a 16-byte number.

g^X generation

- For C-users, g^X is generated in a central server (C-backend server), which knows the sequence of secret values X . The C-backend server maintains a IDTable associated with these secret values (this can be $X = \text{hash}(\text{ID}, \text{epoch})$, etc), with a structure similar to the one used in ROBERT.
- For D-users, g^X is generated in the devices, which keep the sequence of secret values X .

Beacon processing

All users retain the g^X values received, as well as the RSSI of the received beacon.

Users with tests COVID+

C-user is tested COVID+

C-devices choose a secret random number W (which is not known to server C), and insert the list of tuples $[(\text{hash}(g^XW), g^W, \text{RSSI})]$ into a global public list, so that the identity of the device remains anonymous.

D-user is tested COVID+

D-devices choose a secret random number W , different for each received g^X , and insert the list of tuples $[(\text{hash}(g^XW), g^W, \text{RSSI})]$ into a global public list, so that the identity of the device remains anonymous.

The tuples are kept in the public list for a limited period of time, e.g. 1 day. The global list can be organized into geographic areas, in order to reduce the download traffic for D-devices; see section "Support for Roaming".

Exposure notification

C-Users

The C-backend server periodically reads the public list $[(\text{hash}(g^XW), g^W, \text{RSSI})]$. Then it looks for intersections of $\text{hash}((g^W)^X)$, for the X values stored in the IDTable, with the hash values of the list $[(\text{hash}(g^XW), g^W, \text{RSSI})]$.

For C-users, who use the same W value for all published tuples in the list, the IDTable can obtain a time series of contacts for the risk scoring algorithm (because the user with test COVID+ publishes a constant g^W value). If the user who tested COVID+ is a D-user, the time series information is lost.

Note, however, that the C-backend server does not know the identity of the users (C or D) who tested COVID+, because the W value is kept secret in the device.

The C-devices periodically connect to the C-backend server for receiving an Exposure Notification.

D-Users

D-devices periodically read the public list $[(\text{hash}(g^XW), g^W, \text{RSSI})]$. Then, they look for the intersections of the $\text{hash}((g^W)^X)$, for the X values stored in the device, with the hash values in the list $[(\text{hash}(g^XW), g^W, \text{RSSI})]$.

A Risk Scoring is then run in the D-device that decides whether the D-user is notified as in risk of exposure.

We discuss later how to reduce the number of computations in the case of D-users by using Country Codes.

Risk scoring

The proximity of the contact can be estimated from the RSSI value. The length of the contact is obtained from the times when the X value was used. Centralized backend servers can use time series in their risk scoring algorithms, since the W value in the C-devices is the same for all declared tuples $(\text{hash}(g^XW), g^W, \text{RSSI})$.

Support to roaming

The devices could add a country code (e.g. CC = "Spain", "France", etc.) to each tuple in the published list: $[(\text{hash}(g^XW), g^W, \text{RSSI}, \text{CC})]$.

This country code corresponds to the place where the g^X was received. Users should specify the country where they are located to the application. If no location is given, a special code can be used ("Schengen"). Since the identity of users is kept anonymous, users do not reveal publicly their location.

The C-backend server knows the location where the interaction of the user who transmitted g^X took place, but this is already the situation in ROBERT in the case of federated backend servers.

Note that for D-users, their location remains private.

C-Users

C-backend servers must check the intersections for the entire list published in the Schengen area. This calculation is indeed long $(14 \times 100 \times \text{length}([(\text{hash}(g^XW), g^W, \text{RSSI}, \text{CC})]))$ but it is performed twice a day, and should not be a major problem for a backend server.

D-Users

D-devices should only evaluate the $\text{hash}((g^W)^X)$ for elements with a CC that corresponds to one of the zones visited in the last 14 days (or the cases $\text{CC}=\text{"Schengen"}$).

Defences again other attacks

Devices add to the list two more fields: $\text{hash}(g^X + \text{epoch})$, $\text{hash}(g^X + \text{MAC})^4$, where epoch is the time of reception of g^X , and MAC is the MAC address of the beacon that contained g^X . These fields need to be checked only in case of intersections.

Consequences of breaking Diffie-Hellman

If an attacker breaks DH (i.e. gets W from g^W), she could check if a g^X that has been eavesdropped matches the $\text{hash}(g^XW)$ value. This would lead to the conclusion that the user who transmitted g^X was close to a user who reported a COVID+ test. This vulnerability is inherent in all digital proximity tracing protocols, and the attacker can obtain this information in a much simpler way.

Possible implementation issues

- The current Gapple EN API would not support this mode of operation.
- Is 16-byte Diffie-Hellman too weak?. We do not think so, as the information that an attacker can obtain from breaking DH is not worth the effort. However, in case this is considered a risk, using 32 bytes DH and the corresponding consequences on beacon transmissions should be considered.

Privacy properties

D-users avoid re-identification attacks of distributed protocols such as DP3T. C-users share less information with the C-backend server, as users are keep anonymous when reporting a test COVID+.

Acknowledgement

The author would like to thank Leonardo Bautista-Gómez, Pedro Martín-Jurado, and Pablo Rodríguez-Rodríguez for useful discussions on the properties of this protocol. The fact that IDPT can be used as a full digital proximity tracing protocol was clear from the beginning, but the author made no effort to work out the details until Pablo pointed out that this protocol would probably not suffer from the re-identification problems of other distributed protocols. However, any errors in the text are the sole responsibility of the author.

⁴ Is this enough: $\text{hash}(g^X + \text{epoch} + \text{MAC})$?